

Twitter Sentiment Analysis - Deep learning project

This document outlines the iterative process used to develop a robust sentiment analysis model and describes the production-ready infrastructure implemented for high-concurrency inference.

The goal was to classify text sentiment accurately and deploy the model in an environment capable of handling high user traffic, specifically benchmarking against **500 concurrent users**.

1. Model Development Iterations

The project began by exploring various classification scopes and fine-tuning strategies using Hugging Face's Transformers library, primarily focusing on `TFBertForSequenceClassification` and `TFDistilBertForSequenceClassification`.

Attempt	Classification Scope	Model & Strategy	Outcome & Challenge
1	4 Classes (Positive, Negative, Neutral, Irrelevant)	<code>TFBertForSequenceClassification</code> + Class Weights	High complexity due to large parameter count; infeasible for rapid iteration and deployment resource constraints. Dropped due to the size of the BERT model.
2	4 Classes (Positive, Negative, Neutral, Irrelevant)	<code>TFDistilBertForSequenceClassification</code> + Full Fine-tuning + Class Weights	Training was excessively slow due to full fine-tuning, despite using the smaller DistilBERT model.
3	2 Classes (Positive, Negative)	<code>TFDistilBertForSequenceClassification</code> + Full Fine-tuning	Accuracy dropped significantly (<code>< 0.5</code>), indicating severe overfitting. Removing less-represented classes (Neutral, Irrelevant) did not immediately solve the generalization issue with full fine-tuning.

4	2 Classes (Positive, Negative)	TFDistilBertForSequenceCl assification + Frozen Weights	Introduced a frozen weight strategy (transfer learning) with a moderate learning rate (). Model continued to overfit, yielding low test accuracy.
5 (Final)	2 Classes (Positive, Negative)	TFDistilBertForSequenceCl assification + Frozen Weights + Regularization	Successfully mitigated overfitting using a multi-pronged approach, achieving the required performance target.

2. Final Model Configuration (Approach 5)

The final, successful approach focused on preventing overfitting in the pre-trained DistilBERT architecture, resulting in a high-accuracy, deployable model:

Parameter	Value / Description
Base Model	TFDistilBertForSequenceClassification
Classification Scope	2 Labels (Positive, Negative)
Fine-tuning Strategy	Frozen Weights Training (Only top layers are trainable)
Learning Rate	Reduced to 1e-4 (from 5e-4)
Regularization	Dropout Layer with a rate of 0.4
Optimization	Adam Optimizer
Weight Decay	L2 Regularization with weight decay 0.0001
Training Controls	Early Stopping and Model Checkpoint (to save only the best performing model)
Test Accuracy	Achieved 80.97%
Precision for class 0 and 1	0.8056 , 0.8146

Recall for class 0 and 1	0.8356, 0.7817
F1-score for class 0 and 1	0.8203, 0.7978

3. Deployment, Optimization, and Load Testing

The trained model was containerized and deployed using a scalable architecture designed for high availability and low latency inference.

Prediction API (FastAPI)

- **Framework:** FastAPI was chosen for its modern, asynchronous nature, enabling high performance.
- **Optimization:**
 - **Single Model Load:** The model is loaded only **once** into memory upon application startup, eliminating disk I/O latency for every prediction request.
 - **Threading for Inference:** Standard Python threading was implemented to prevent resource allocation conflicts during model inference, ensuring non-blocking execution and smoother parallel processing.

Deployed API Endpoint

The prediction API is deployed and accessible at the following URL:

<https://dac7a9adfa2824d729f097cdc2dd2cf65.clg07azjl.paperspacegradient.com/predict>

Sample Request Payload (POST):

```
{
  "text": "I am really frustrated by your service!"
}
```

Sample Response:

```
{
  "sentiment": "Negative",
  "confidence_score": 0.7034
}
```

Dockerization and Scalability

- **Environment:** The entire application (FastAPI + Model) was containerized using **Docker**.

- **Hardware:** Deployed on a cloud machine with **8 CPU cores, 30 GB RAM, GPU support,** and **auto-scaling** capability.
- **Concurrency Handling (Docker Compose):** A docker-compose setup was implemented with **7 worker containers** running the FastAPI service, allowing the prediction API to handle high concurrent load by leveraging multiple cores.

Key Docker Commands:

Command	Purpose
<code>docker pull afrahthahir/sentiment-api:latest</code>	Pulls the docker image from the dockerhub
<code>docker compose up -d --scale worker=7</code>	Builds, creates, and starts the 7 worker containers specified in <code>docker-compose.yml</code> in detached mode.
<code>docker-compose down</code>	Stops and removes the containers, networks, and volumes defined in the <code>docker-compose.yml</code> file.


Load Testing (Locust)

- **Tool: Locust** was used to perform performance and load testing against the deployed API.
- **Benchmark Goal:** The system was successfully tested to ensure stable performance and responsiveness when subjected to a load of **500 concurrent users**.


Key Locust Commands:

Command	Purpose
Once the docker compose up -d --scale worker=7 is executed, Go to localhost port 8089 . You will see the locust UI	Starts the Locust testing framework with the multiple workers to handle the distributed api load in your localhost port 8089
Locust Web UI Settings	Configuration used to simulate the target load: <ul style="list-style-type: none"> - Number of users: 500 - Spawn rate: 50 users/second

Total Requests per Second



Number of Users




When the number of concurrent users is 500, the Request Per Second is 15.5, and the Request failed is 0. This shows that the API is scaled to handle 500 concurrent users.

Speech-to-Text Service Implementation for Real-Time Application - R&D Skills:

I have conducted an extensive search for both well-established commercial services and promising open-source solutions suitable for real-time, low-latency applications.

STT Models Comparison Report

My observations and comparisons are recorded in this following report.

 STT Models Comparison Report

Criteria for Evaluation

The following custom criteria were developed to evaluate candidates based on the application's non-negotiable requirement for a robust, real-time experience:

1. **Accuracy (Word Error Rate - WER):** The paramount metric. The solution must provide industry-leading accuracy in various acoustic environments.
2. **Real-Time Latency ($\leq 500\text{ms}$):** Must support a streaming API with end-to-end latency below 500ms to ensure a fluid, conversational user experience.
3. **Cost (Commercial Only):** Evaluation of pricing structures (per minute/second) and tiers against the total cost of ownership (TCO) compared to the effort of managing an open-source solution.